

Image denoising with deep learning

Course Maths and Data sciences

Master mathématiques appliquées, statistique - 2ème année
Parcours Data science.

Frédéric Richard

AMU, 2025

Outline

① Image denoising

② Neural networks

- Overview
- Learning issues
- Evolution of CNN

③ Denoising CNN

- Auto-encoder
- DnCNN

Image Denoising Problem

- Goal: recover a clean image x from a noisy observation y .
- For instance, recover x when corrupted by an additive noise n :

$$y = x + n.$$

- Bayesian viewpoint:

- **Degradation model** $p(x|y)$. For instance, if $n \simeq \mathcal{N}(0, \sigma^2 \text{Id})$ and is independent of x , then

$$p(x|y) \propto \exp\left(-\frac{1}{2\sigma^2}|x - y|^2\right).$$

- **Prior distribution** $p(x)$ of clean images. For instance, assuming some regularity of image, we can use a Tikhonov regularisation:

$$p(x) \propto \exp(-|Dx|^2)$$

on derivatives Dx of x .

- Using Bayes'rule, we obtain the **posterior distribution**:

$$p(x|y) = \frac{p(y|x)p(x)}{p(y)}.$$

Denoising

- The denoiser can be found as the Maximum A Posteriori (MAP) estimator:

$$\begin{aligned}\hat{x} &= \arg \max_x p(x|y) = \arg \max_x p(y|x)p(x) \\ &= \arg \min_x (-\log(p(y|x)) - \log p(x)).\end{aligned}$$

- Alternately, the denoiser can be searched by minimizing the mean-square error (MSE)

$$\mathbb{E}(|y - x|^2),$$

leading to the Minimum MSE (MMSE):

$$\tilde{x} = \mathbb{E}(x|y) = \int xp(x|y)dx.$$

- Both, MAP and MMSE are reliant on the posterior distribution.

Evolution of research

- A key feature is the modeling of $p(x)$, which has driven the literature of image denoising.
- Classical era (around 1970-2010).
 - Design of $p(x)$, mostly in the form of a Gibbs distribution

$$p(x) \propto \exp(-\rho(x)),$$

where ρ is an energy function.

- Design and study appropriate algorithms to find the MAP depending on the chosen distribution.
- AI revolution (starting around 2012 to nowadays).
- Another approach:
 - Gather a large dataset of clean images $(x_k)_k$ and create noisy image (y_k) from these images,
 - Design a neural network taking y as input and giving an estimate of x as output.
 - Train the neural network by minimising a loss defining the distance between the clean image x and its estimate \hat{x} .

Further reading: A survey by M. Elad et al.

Architecture of a sequential neural network

- Series of U layers : for $u = 1, \dots, U$,

$$f^{(u)} = L^{(u)}(f^{(u-1)}; w^{(u)}),$$

with

- $L^{(u)}$: the layer at level u .
- $f^{(u-1)}$: input signal at level u ,
- $f^{(u)}$: output signal at level u ,
- $w^{(u)}$: parameters at level u .
- Classical components of a layers:
 - Dense layer $g^{(u)} = W^{(u)} f^{(u)}$: linear link between input and output signals.
 - Convolutional layer $g^{(u)} = w^{(u)} * f^{(u)}$: spatially-invariant linear transform.
 - Activation function (\tanh , sigmoid, ReLU, etc.): non-linear representation.
 - Max or Average pooling: local summary of spatial information.
 - Down- (resp. up-) sampling: reduction (resp. increase) the signal size.

Early CNNs: LeNet (1990s)

One of the first CNNs, applied to handwritten digit recognition.




Figure: Image from Medium paper of Abishek Jain.

Typical block of a convolutional layer :

- Convolutions with several kernels: $g^{(u)} = S^{(u)}(f^{(u-1)}; w^{(u)})$
- Activation with a non-linear function φ : $h^{(u)} = \varphi(g^{(u)})$
- Pooling (with down-sampling): $f^{(u)} = P^{(u)}(h^{(u)})$.

Learning a neural net.

- Learning set: $(x_i, y_i)_{i=1}^n$.
- Predict y_i from x_i with the output of the last layer $\hat{y}_i(w)$.
- $\hat{y}_i(w)$ is defined by recursion:
 $f_i^{(0)} = x_i$ and for all $u = 1, \dots, U$, $f_i^{(u)} = L^{(u)}(f_i^{(u-1)}, w^{(u)})$.
- $w = (w^{(u)}, u = 1, \dots, U)$ is the set of model parameters.
- Minimisation of the empirical risk

$$E_n(w) = \frac{1}{n} \sum_{i=1}^n \ell(y_i, \hat{y}_i(w)),$$

defined for some loss function ℓ (mean square error, binary-cross entropy, etc.)

Optimization

- Minimization by stochastic gradient algorithm and variants.
- involve computing gradients of the loss function.
- Back-propagation: recursive expression of gradients through layers:
- Typical example:

$$E_n(w) = \frac{1}{n} \sum_{i=1}^n \ell(y_i, f_i^{(U)}) \quad \text{with} \quad \begin{cases} g_i^{(u)} &= S^{(u)}(f_i^{(u-1)}; w^{(u)}), \\ f_i^{(u)} &= \varphi^{(u)}(g_i^{(u)}). \end{cases}$$


$$\frac{dE_n}{dw^{(v)}}(w) = \frac{1}{n} \sum_{i=1}^n \frac{d\ell}{dz} \left(y_i, \frac{df_i^{(L)}}{dw^{(v)}} \right)$$

$$\frac{df_i^{(u)}}{dw^{(v)}} = d\varphi^{(u)}(g_i^{(u)}) \frac{dg_i^{(u)}}{dw^{(v)}}.$$


Vanishing gradient and activation functions

$$u = 1, \dots, U, \quad \left| \frac{df_i^{(u)}}{dw^{(v)}} \right| \leq \left| d\varphi^{(u)}(g_i^{(u)}) \right| \left| \frac{dg_i^{(u)}}{dw^{(v)}} \right|.$$


- Differential norm is not of the same order through the layers,
- Activation functions $\varphi^{(u)}$ may saturate and make the differential vanish ($|d\varphi^{(u)}(g_i^{(u)})| \simeq 0$).
- Consequence: slow convergence on last layers.



Logistic
 $\varphi(t) = \frac{1}{1+e^{-t}}$



ReLU
 $\varphi(t) = \max(0, t)$



leaky ReLU(α)
 $\varphi(t) = \max(\alpha t, t)$

Internal covariate shift and batch normalization

- Issue: statistics of layer output vary drastically from a layer to another, resulting in instabilities.
- Solution (Ioffe et Szegedy, 2015) : centering and normalizing the output with batch normalization:
- Example:
 - Output of the layer u :

$$g_i^{(u)} = S^{(u)}(f_i^{(u-1)}; w^{(u)}).$$

- Estimation of the mean and variance of $g^{(u)}$ from a batch B :


$$\hat{\mu}_B^{(u)} = \frac{1}{|B|} \sum_{i \in B} g_i^{(u)} \text{ et } \hat{\sigma}_B^{(u)} = \frac{1}{|B|} \sum_{i \in B} (g_i^{(u)} - \hat{\mu}_B^{(u)})^2.$$

- Centering and normalizing the output.

$$\tilde{g}_i^{(u)} = \frac{1}{\sqrt{\hat{\sigma}_B^{(u)}} + \varepsilon} (g_i^{(u)} - \hat{\mu}_B^{(u)}).$$

Overfitting

- Issue: Although the model performs well on the learning set, it gives poor results on test data (not seen during learning).
- Factors:
 - Model is too complex.
 - Learning set is too small or not enough representative of all data.
- Early diagnostic during the training step with a validation dataset.



Some solutions to over-fitting.

- Penalization of the optimisation problem, implicitly reducing model complexity.
- **Dropout** (Hinton, 2012 et Srivastava et al. 2014):
 - Cancel randomly and temporarily the evolution of weights during the learning,
 - Attenuate specialization of weights on learning data.
- **Data augmentation**: increase the learning set using data transforms (translation, rotation, contrast changes, etc.).
- Fine tuning: use an already learned neural network and modify only a few of its parameters to adapt it to another context.u classifieur.

From LeNet to AlexNet

- **LeNet**: one of the first CNNs, applied to handwritten digit recognition.
- Limitations:
 - Vanishing gradients in deep networks.
 - Sigmoid / tanh activations \Rightarrow slow convergence, computational cost.
 - Hardware at the time could not support large CNNs.
- AlexNet: Winner of ImageNet Large Scale Visual Recognition Challenge (ILSVRC, 2012).
- Key innovations:
 - Large set of data for training **ImageNet**.
 - **GPU** acceleration for training (Cuda programming, Nvidia).
 - **ReLU** activations instead of sigmoid and tanh \Rightarrow faster training.
 - Local response normalization \Rightarrow stability.
 - **Dropout** layers to attenuate over-fitting


CNN of ILSVRC

- **AlexNet** - milestone in deep learning.
 - Limitation: very large kernels \Rightarrow high memory usage.
- **VGG** (Winner of ILSVRC 2014)
 - Used many **stacked small** 3×3 **filters** \Rightarrow parameter reduction.
 - Deeper architecture improved accuracy.
 - Drawback: very deep \Rightarrow risk of vanishing gradients.
- **GoogLeNet** (2014)
 - **Inception module**.
 - Replaced large filters with multiple smaller ones to reduce parameters.
 - Increased network **width** instead of only depth.
 - Drawback: wide networks may suffer from overfitting.
- **ResNet** (Winner of ILSVRC 2015).
 - Introduced **residual connections**: $y = F(x) + x$.
 - Allowed training of very deep networks without vanishing gradients.
 - Foundation for many modern architectures.

Further reading: Medium paper on AlexNet, VGGNet, ResNet, and Inception.

Autoencoder Basics

- Encoder compresses the input image into a low-dimensional **latent representation**.
- Decoder reconstructs the image from this latent representation.
- For denoising: train the autoencoder to reconstruct the clean image from a noisy input.



Convolutional auto-encoder for image denoising:

- Encoder:** convolution + pooling + downsampling layers.
- Latent space:** compressed representation of images.
- Decoder:** upsampling + transposed convolution layers.
- Output:** denoised image \hat{y} .

Auto-encoder for image denoising

- Training Objective:

- Dataset: pairs of noisy images x_i and clean images y_i .
- Loss function: Mean Squared Error (MSE)

$$\mathcal{L}(\theta) = \frac{1}{N} \sum_{i=1}^N \|f_\theta(y_i) - x_i\|^2$$

- Alternative: Mean Absolute Error (L1), SSIM-based loss, perceptual loss.
- Advantages:
 - Simple architecture.
 - Forms the basis for more advanced denoising models (DnCNN, U-Net, VAEs, diffusion).
- Difficulty to recover the image details.

U-Net Architecture

- Originally proposed by Ronneberger et al. (2017) for biomedical segmentation.
- Encoder–decoder architecture with **skip connections**.
- Captures both global context and fine local details.
- **Contracting part:**
convolution + downsampling
(captures context).
- **Bottleneck:**
deepest layer with high-level features.
- **Expansive part:**
upsampling + convolution
(restores resolution).
- **Skip connections:**
transfer fine details.




Figure: source: (Ronneberger et al., 2015)


DnCNN Overview

- Overview:

- Proposed by Zhang et al. (2017).
- Deep CNN specifically designed for image denoising.
- Uses **residual learning**: instead of directly predicting x , the network predicts the noise \hat{n} .
- Output: $\hat{x} = y - \hat{n}$.

- Architecture:

- Input: noisy image $x \in \mathbb{R}^{H \times W \times C}$.
- First layer: Conv + ReLU.
- Middle layers: several Conv + BatchNorm + ReLU blocks.
- Last layer: Conv (predicts residual noise).



Training Objective

- Training objective

- Training data: pairs (x, y) of noisy/clean images.
- Loss function: Mean Squared Error (MSE) between predicted noise \hat{n} and true noise n :

$$\mathcal{L}(\theta) = \frac{1}{N} \sum_{i=1}^N \| \hat{n}_i - (y_i - x_i) \|^2$$

- Equivalent to a skip connection with a classical MSE.

- Advantages:

- Residual learning simplifies optimization.
- Batch Normalization stabilizes training.
- Deeper CNN captures spatial context for stronger denoising.
- Can generalize to different noise levels (blind denoising).