

Large Language Models in Data Science

Week 5: Retrieval-Augmented Generation (RAG)

Sebastian Mueller

Aix-Marseille Université

2025-2026



Session Overview

Lecture (1h)

1. From LLMs to RAG
2. Retrieval and search basics
3. Augmenting prompts with context
4. Lexical, metadata, and embedding search
5. Putting it together: RAG architectures
6. Example: Marseille & data science assistant

Lab (2h)

- ▶ Ingest a small corpus about Marseille & AMU data science
- ▶ Build an embedding index with metadata
- ▶ Implement a simple RAG chain
- ▶ Compare answers with and without retrieval
- ▶ Optional: add citations & reranking


Where We Are in the Course

- ▶ Week 1: Tokens, embeddings, transformer architecture.
- ▶ Week 2: Using pretrained models via Hugging Face.
- ▶ Week 3: Effective LLM use for coding, research, ideation.
- ▶ Week 4: Text classification and intent routing for an AMU chatbot.
- ▶ **Week 5: RAG** combine all ingredients to build assistants grounded in your own data.

Why Do We Need RAG?

- ▶ LLMs are trained on huge but *fixed* corpora; they do not know your internal documents or freshest data.
- ▶ Direct prompting relies on the model “remembering” facts from pretraining, which leads to hallucinations and outdated answers.
- ▶ Many data science tasks need answers grounded in:
 - ▶ institutional documents (course catalog, policies),
 - ▶ project reports, notebooks, dashboards,
 - ▶ domain-specific knowledge bases.
- ▶ RAG attaches a *retrieval system* to the LLM so answers can depend on external, updatable data.

From Vanilla LLM to RAG



From Vanilla LLM to RAG

- ▶ Mathematically we move from $P(\text{answer} \mid \text{question})$ to
 $P(\text{answer} \mid \text{question, retrieved context}).$
- ▶ The only thing the LLM sees is the *prompt*; RAG works by **augmenting the prompt** with relevant context.

Core Components of a RAG System

- ▶ **Document store:** where raw texts live (internet, files, database, data lake, etc.).
- ▶ **Chunking:** split documents into passages that are short enough to fit in the context window.
- ▶ **Metadata:** structured fields such as title, source, date, language, tags.
- ▶ **Index:** data structures to support fast search (inverted index, vector index, or both).
- ▶ **Retriever:** given a query, returns relevant chunks (possibly with scores).
- ▶ **LLM + prompt template:** turns query + retrieved chunks into an answer.

Documents, Chunks, and Metadata

- ▶ **Document:** a logically complete piece of text (course syllabus, FAQ page, PDF report).
- ▶ **Chunk:** a slice of a document (e.g., a few paragraphs) used as the atomic retrieval unit.
- ▶ **Metadata examples for AMU / Marseille:**
 - ▶ source_type (“course_catalog”, “housing_info”, “lab_sheet”).
 - ▶ campus (“Luminy”, “Saint-Charles”, “Aix”), city (“Marseille”, “Aix-en-Provence”).
 - ▶ program (“MSc Data Science”), year, language.

Retrieval Basics: Three Signals

- ▶ **Lexical / keyword search**
 - ▶ Uses exact or fuzzy matches of words (e.g., BM25).
 - ▶ Very precise when user vocabulary matches the documents.
- ▶ **Embedding (semantic) search**
 - ▶ Encode query and chunks into vectors; retrieve nearest neighbors.
 - ▶ Captures semantics beyond exact words (“student card” \approx “carte étudiant”).
- ▶ **Metadata filtering**
 - ▶ Restrict the candidate set using structured filters (e.g., campus = “Marseille”).
 - ▶ Cheap and interpretable; in our lab pipeline we apply it as a *final refinement* after lexical and embedding search.
 - ▶ In large production systems you might also use metadata as a coarse pre-filter before more expensive retrieval.

From Bag-of-Words to BM25

- ▶ **Bag-of-words (BoW):**
 - ▶ represent each document as a vector of word counts,
 - ▶ simple and fast, but treats all terms and all documents equally.
- ▶ **Term frequency (TF):** words that repeat many times in a document should matter more.
- ▶ **Inverse document frequency (IDF):** rare words across the corpus carry more signal than very common ones.
- ▶ Naive TF or TF-IDF:
 - ▶ can over-favor very long documents (more opportunities to match),
 - ▶ can over-reward repeated words without saturation.
- ▶ BM25 refines this idea to give a stronger, length-aware lexical baseline.

BM25: Tunable Keyword Scoring

- ▶ **Definition:** BM25 scores a document d for a query term t roughly as:

$$\text{score}(t, d) \propto \text{IDF}(t) \frac{(k_1 + 1) \text{TF}(t, d)}{\text{TF}(t, d) + k_1 \left(1 - b + b \cdot \frac{\text{document length}}{\text{average document length}}\right)}.$$

Best Matching 25 was named as the 25th variant in a series of scoring functions proposed by its creators.


- ▶ **Key intuitions:**

- ▶ *Term frequency saturation:* repeating a word helps, but with diminishing returns.
- ▶ *Document length normalization:* longer documents are penalized.
- ▶ *IDF:* rare terms across the corpus contribute more to the score.

- ▶ **Tunable parameters:**

- ▶ k_1 controls how fast the TF effect saturates,
- ▶ b controls how strongly document length is normalized.
- ▶ In practice: a well-tuned BM25 is often the *best starting point* for keyword-based retrieval in RAG systems.


Combining Keywords, and Embeddings



Combining Keywords, and Embeddings

- ▶ Typical RAG uses a **hybrid retriever**:
 - ▶ run keyword search and embedding search,
 - ▶ combine candidates by intersection / union and rerank,
 - ▶ then apply metadata filters as a final refinement step (as in our lab).
- ▶ Choice of combination controls the trade-off between recall and precision.

From Retrieval to Augmented Prompt



- ▶ RAG is not a new model; it is a *prompting strategy* powered by a retriever.
- ▶ Quality hinges on:
 - ▶ what you retrieve (indexing, chunking, retriever),
 - ▶ how you format it in the prompt (templates, instructions, citations).

Prompt Template for Q&A

- ▶ Typical template (simplified):

You are a helpful assistant for data science students in Marseille.
Use only the information in the CONTEXT to answer.
If the answer is not in the CONTEXT, say you don't know.

CONTEXT:

```
{{ retrieved_chunks }}
```

QUESTION:

```
{{ user_question }}
```

- ▶ The retrieved chunks become part of the prompt; they *augment* the original question.
- ▶ Clear instructions reduce hallucinations and encourage citing the provided context.

Context Size and Ranking

- ▶ Context window is limited: we cannot paste the entire knowledge base into every prompt.
- ▶ We must choose:
 - ▶ **Top- k** : how many chunks to include?
 - ▶ **Max tokens**: truncate text to stay within the limit.
 - ▶ **Ranking strategy**: score by similarity, recency, or custom signals (e.g., student vs. admin).
- ▶ Trade-off:
 - ▶ too few chunks \Rightarrow model misses key information;
 - ▶ too many chunks \Rightarrow context becomes noisy and harder to use.

Use Case: AMU / Marseille Data Science Assistant

- ▶ Build an assistant for:
 - ▶ the MSc Data Science programme in Marseille,
 - ▶ practical life in the city for students.
- ▶ Example questions:
 - ▶ “Where are the data science lectures in Marseille usually held?”
 - ▶ “How can I access the computer labs on Sundays?”
 - ▶ “What are good quiet places to study data science in Marseille?”
- ▶ We want answers grounded in:
 - ▶ official programme pages,
 - ▶ campus maps and schedules,
 - ▶ curated local tips in our own documents.

Step 1: Ingest & Chunk

- ▶ Collect source documents:
 - ▶ programme descriptions (PDFs, HTML),
 - ▶ campus information and opening hours,
 - ▶ local guides curated by teaching staff.
- ▶ Normalize text (encoding, language detection, basic cleaning).
- ▶ Chunk into passages:
 - ▶ target e.g. 200–500 tokens per chunk,
 - ▶ keep semantic coherence (sections, headings).
- ▶ Attach metadata:
 - ▶ campus=Marseille, program=Data Science, source_url, section_title, ...

Step 2: Index & Retrieve

- ▶ Compute embeddings for each chunk using a multilingual sentence model.
- ▶ Store vectors in a vector index (e.g., FAISS, ScaNN, or a hosted service).
- ▶ At query time:
 1. Parse the user question (language, intent).
 2. Run hybrid search (embeddings + keywords) over all candidate chunks.
 3. Apply metadata filters as a final refinement (e.g., campus=Marseille, program=Data Science).
 4. Return top- k chunks with scores after filtering.
- ▶ This “search” stage is independent from the LLM and can be tuned separately (retriever, metadata filters, and ranking strategy).

Step 3: Build the Augmented Prompt - Template in code (pseudo-Python)

```
chunks = retrieve_hybrid(query) # BM25 + embeddings
chunks = [ # Metadata-based refinement applied after retrieval
    c for c in chunks
    if c["campus"] == "Marseille"
    and c["program"] == "Data Science"
]
context = "\n\n".join(chunk["text"] for chunk in chunks)
prompt = f"""
You are an assistant for data science students in Marseille. Answer the question using only the CONTEXT below. If the answer is not in the CONTEXT, say you don't know.
CONTEXT:
{context}
QUESTION:
{query}"""
```

Step 4: Answer, Evaluate, Iterate

- ▶ Send the augmented prompt to the chosen LLM (OpenAI, local model, etc.).
- ▶ Log:
 - ▶ user question,
 - ▶ retrieved chunks and scores,
 - ▶ final answer and model parameters.
- ▶ Evaluate on a small test set:
 - ▶ correctness (does the answer match ground truth?),
 - ▶ grounding (can we see the supporting chunks?),
 - ▶ robustness (different phrasings, languages, levels of detail).
- ▶ Iteratively tune:
 - ▶ retriever (chunking, embeddings, ranking),
 - ▶ prompt (instructions, formatting, citations).

Precision and Recall Refresher

- ▶ **Precision:** among the retrieved items, how many are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}.$$

- ▶ **Recall:** among all relevant items in the corpus, how many did we retrieve?

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}.$$

- ▶ Example (retrieval for a RAG system):

- ▶ there are 20 truly relevant chunks in the corpus,
- ▶ your retriever returns 10 chunks, 8 of which are relevant.

Then:

$$\text{Precision} = \frac{8}{10} = 0.8, \quad \text{Recall} = \frac{8}{20} = 0.4.$$

- ▶ For RAG, we usually prefer *high recall* at the retrieval stage, then let the LLM focus on precision.

RAG vs Fine-Tuning

- ▶ **Fine-tuning:**
 - ▶ adapts model weights,
 - ▶ good for new behaviors or styles,
 - ▶ expensive to update, may still hallucinate facts.
- ▶ **RAG:**
 - ▶ keeps base model fixed,
 - ▶ plugs in external, updatable knowledge via retrieval,
 - ▶ easier to iterate on (update data, retriever, or prompts).
- ▶ In practice:
 - ▶ start with RAG for most data-centric use cases,
 - ▶ consider fine-tuning on top of a good RAG pipeline if behavior still needs shaping.

Common Pitfalls

- ▶ **Hallucinated citations:** model invents sources or mixes documents.
- ▶ **Shallow retrieval:** top- k chunks are off-topic or redundant.
- ▶ **Stale data:** document store not updated as the real world evolves.
- ▶ **Leaky prompts:** model uses prior knowledge instead of the provided context.
- ▶ **Evaluation gap:** no clear test set or feedback loop.

- ▶ Good engineering practice:
 - ▶ log everything,
 - ▶ add human review for early deployments,
 - ▶ maintain small, curated evaluation sets.

How the Pieces Fit Together

- ▶ Tokens, embeddings, and transformers (Weeks 1–2) explain how LLMs represent and process text.
- ▶ Effective prompting (Week 3) gives us control over model behavior.
- ▶ Classification (Week 4) is one example of using embeddings and LLMs for structured decisions.
- ▶ **RAG (Week 5)** combines all of these:
 - ▶ embeddings for retrieval,
 - ▶ prompts for grounding and style,
 - ▶ evaluation for reliability.
- ▶ This is the main pattern behind many modern data science assistants.

Looking Ahead

- ▶ In the lab you will:
 - ▶ build a minimal RAG pipeline over a small Marseille / AMU corpus,
 - ▶ experiment with keyword vs embedding vs hybrid retrieval,
 - ▶ observe how prompt design affects grounding.
- ▶ In the hackathon you can reuse these ideas:
 - ▶ many impactful projects are just: “a good RAG system + a good UI”.